



東南大學
SOUTHEAST UNIVERSITY

OPERATING SYSTEM CONCEPTS

.....

Chapter 13. I/O Systems

A/Prof. Kai Dong

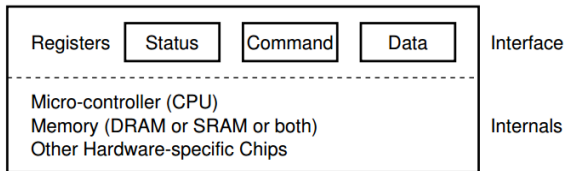
Contents





A Canonical Device

- A device has two important components.
 - The first is the hardware interface it presents to the rest of the system
 - The second part of any device is its internal structure, which is implementation specific.



- Three registers:
 - Status — to be read to see the current status of the device;
 - Command — to tell the device to perform a certain task;
 - Data — to pass data to , or get data from the device.



A Canonical Protocol

- Polling is inefficient

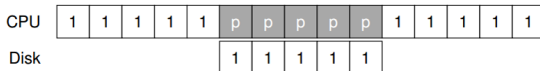
```
1 While (STATUS == BUSY)
2     ;           // wait until device is not busy
3 Write data to DATA register
4 Write command to COMMAND register
5     (Doing so starts the device and executes the command)
6 While (STATUS == BUSY)
7     ;           // wait until device is done with your request
```



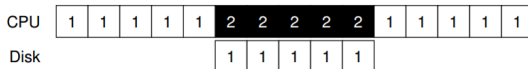
Interrupts

- Interrupt-driven
- Instead of polling the device repeatedly, the OS can issue a request, put the calling process to sleep, and context switch to another task. When the device is finally finished with the operation, it will raise a hardware interrupt, causing the CPU to jump into the OS at a pre-determined interrupt service routine (ISR) or more simply an interrupt handler.

– Without interrupts



– Interrupt-driven





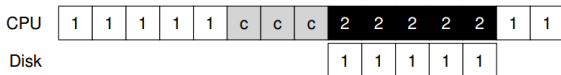
Interrupts

- Take the device speed into account
 - If a device is fast, it may be best to poll; if it is slow, interrupts are best.
 - If the speed of the device is not known, or sometimes fast and sometimes slow, it may be best to use a hybrid that polls for a little while and then, if the device is not yet finished, uses interrupts.
- Livelocks
 - In networks, when a huge stream of incoming packets each generate an interrupt, it is possible for the OS to livelock, that is, find itself only processing interrupts and never allowing a user-level process to run and actually service the requests.
 - Slashdot effect
 - Coalescing: multiple interrupts can be coalesced into a single interrupt delivery

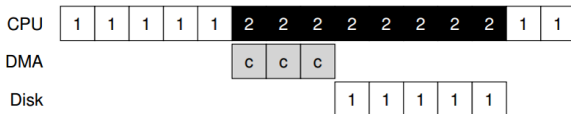


Direct Memory Access

- One remaining problem:
- Suppose a transfer of a large chunk of data to a device
 - The CPU is overburdened



- Direct Memory Access (DMA).
 - A DMA engine is essentially a very specific device within a system that can orchestrate transfers between devices and main memory without much CPU intervention.





Direct Memory Access

- How to communicate with devices?
 - I/O instructions.
 - » E.g., the in and out instructions in x86.
 - » Such instructions are usually privileged.
 - Memory mapped I/O.
 - » The hardware makes device registers available as if they were memory locations.



Device Driver

- Device driver
 - A piece of software at the lowest level in the OS which know in detail how a device works.

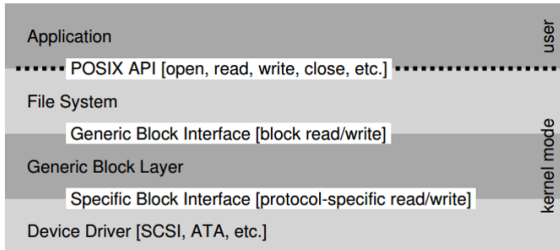


Figure 36.3: The File System Stack